

*Práctica 6: Regresión con AutoEncoders Convolucionales***Denoising**

**7. Describe el mejor modelo que hayas encontrado: balance entre mejor desempeño y menos parámetros. Menciona el tiempo que toma el entrenamiento del modelo seleccionado.**

Layer (type)	Output Shape	Param #
input_17 (InputLayer)	(None, 64, 64, 3)	0
batch_normalization_33 (Batch Normalization)	(None, 64, 64, 3)	12
conv2d_62 (Conv2D)	(None, 64, 64, 8)	224
max_pooling2d_21 (MaxPooling)	(None, 32, 32, 8)	0
conv2d_63 (Conv2D)	(None, 32, 32, 16)	1168
dropout_21 (Dropout)	(None, 32, 32, 16)	0
up_sampling2d_21 (UpSampling)	(None, 64, 64, 16)	0
conv2d_64 (Conv2D)	(None, 64, 64, 8)	1160
batch_normalization_34 (Batch Normalization)	(None, 64, 64, 8)	32
conv2d_65 (Conv2D)	(None, 64, 64, 3)	219
Total params: 2,815		
Trainable params: 2,793		
Non-trainable params: 22		

El tiempo promedio por época es 62s. El mejor desempeño se obtuvo después de 91 épocas.

El tiempo total de entrenamiento fue de 1 hora y media.

Es el mejor modelo porque la pérdida de validación es de 0.0031 y tiene 2,792 parametros.

El modelo tiene 4 capas convolucionales. La primera y la tercera tienen 8 filtros. La segunda incrementa el número de filtros a 16 y finalmente la última tiene 3 filtros.

Las entradas se normalizan utilizando Batch Normalization. A continuación, se disminuyen los datos con la operación max-pooling y después de una convolución con 16 filtros, la operación upsampling hace que vuelvan a su tamaño original.

```
i = Input(x_train_noise.shape[1:])
h = BatchNormalization()(i)

h = Conv2D(8, (3, 3), padding='same', activation='elu')(h)
h = MaxPooling2D((2, 2))(h)

h = Conv2D(16, (3, 3), padding='same', activation='elu')(h)
h = Dropout(0.2)(h)

h = UpSampling2D((2, 2))(h)
h = Conv2D(8, (3, 3), padding='same', activation='elu')(h)

h = BatchNormalization()(h)
o = Conv2D(3, (3, 3), padding='same', activation='sigmoid')(h)

model = Model(inputs=i, outputs=o)
```

## **8. Reporta el proceso de minimización de la pérdida por época, del modelo seleccionado.**

Train on 4000 samples, validate on 1000 samples

Epoch 1/200

4000/4000 [=====] - 73s 18ms/step - loss: 0.0617 - val\_loss: 0.0296

Epoch 2/200

4000/4000 [=====] - 71s 18ms/step - loss: 0.0294 - val\_loss: 0.0294

Epoch 3/200

4000/4000 [=====] - 76s 19ms/step - loss: 0.0215 - val\_loss: 0.0257

Epoch 4/200

4000/4000 [=====] - 68s 17ms/step - loss: 0.0173 - val\_loss: 0.0186

Epoch 5/200

4000/4000 [=====] - 63s 16ms/step - loss: 0.0144 - val\_loss: 0.0142


Epoch 6/200  
4000/4000 [=====] - 62s 16ms/step - loss: 0.0128 - val\_loss: 0.0129  
Epoch 7/200  
4000/4000 [=====] - 62s 15ms/step - loss: 0.0118 - val\_loss: 0.0129  
Epoch 8/200  
4000/4000 [=====] - 62s 16ms/step - loss: 0.0111 - val\_loss: 0.0122  
Epoch 9/200  
4000/4000 [=====] - 62s 16ms/step - loss: 0.0105 - val\_loss: 0.0110  
Epoch 10/200  
4000/4000 [=====] - 61s 15ms/step - loss: 0.0101 - val\_loss: 0.0103  
Epoch 11/200  
4000/4000 [=====] - 61s 15ms/step - loss: 0.0097 - val\_loss: 0.0097  
Epoch 12/200  
4000/4000 [=====] - 61s 15ms/step - loss: 0.0094 - val\_loss: 0.0090  
Epoch 13/200  
4000/4000 [=====] - 61s 15ms/step - loss: 0.0090 - val\_loss: 0.0085  
Epoch 14/200  
4000/4000 [=====] - 61s 15ms/step - loss: 0.0088 - val\_loss: 0.0081  
Epoch 15/200  
4000/4000 [=====] - 63s 16ms/step - loss: 0.0085 - val\_loss: 0.0076  
Epoch 16/200  
4000/4000 [=====] - 61s 15ms/step - loss: 0.0083 - val\_loss: 0.0073  
Epoch 17/200  
4000/4000 [=====] - 62s 15ms/step - loss: 0.0081 - val\_loss: 0.0070  
Epoch 18/200  
4000/4000 [=====] - 62s 15ms/step - loss: 0.0079 - val\_loss: 0.0068  
Epoch 19/200  
4000/4000 [=====] - 62s 16ms/step - loss: 0.0077 - val\_loss: 0.0066  
Epoch 20/200  
4000/4000 [=====] - 62s 16ms/step - loss: 0.0075 - val\_loss: 0.0064  
Epoch 21/200  
4000/4000 [=====] - 62s 16ms/step - loss: 0.0073 - val\_loss: 0.0063  
Epoch 22/200  
4000/4000 [=====] - 63s 16ms/step - loss: 0.0071 - val\_loss: 0.0062  
Epoch 23/200  
4000/4000 [=====] - 62s 16ms/step - loss: 0.0070 - val\_loss: 0.0060  
Epoch 24/200  
4000/4000 [=====] - 61s 15ms/step - loss: 0.0069 - val\_loss: 0.0059  
Epoch 25/200  
4000/4000 [=====] - 61s 15ms/step - loss: 0.0067 - val\_loss: 0.0059  
Epoch 26/200  
4000/4000 [=====] - 61s 15ms/step - loss: 0.0066 - val\_loss: 0.0057  
Epoch 27/200  
4000/4000 [=====] - 62s 16ms/step - loss: 0.0064 - val\_loss: 0.0057

Epoch 28/200  
4000/4000 [=====] - 62s 15ms/step - loss: 0.0063 - val\_loss: 0.0056  
Epoch 29/200  
4000/4000 [=====] - 62s 15ms/step - loss: 0.0062 - val\_loss: 0.0055  
Epoch 30/200  
4000/4000 [=====] - 62s 16ms/step - loss: 0.0061 - val\_loss: 0.0054  
Epoch 31/200  
4000/4000 [=====] - 61s 15ms/step - loss: 0.0060 - val\_loss: 0.0054  
Epoch 32/200  
4000/4000 [=====] - 62s 15ms/step - loss: 0.0059 - val\_loss: 0.0052  
Epoch 33/200  
4000/4000 [=====] - 61s 15ms/step - loss: 0.0058 - val\_loss: 0.0052  
Epoch 34/200  
4000/4000 [=====] - 62s 15ms/step - loss: 0.0057 - val\_loss: 0.0051  
Epoch 35/200  
4000/4000 [=====] - 61s 15ms/step - loss: 0.0056 - val\_loss: 0.0051  
Epoch 36/200  
4000/4000 [=====] - 62s 15ms/step - loss: 0.0055 - val\_loss: 0.0050  
Epoch 37/200  
4000/4000 [=====] - 62s 15ms/step - loss: 0.0054 - val\_loss: 0.0049  
Epoch 38/200  
4000/4000 [=====] - 62s 16ms/step - loss: 0.0054 - val\_loss: 0.0049  
Epoch 39/200  
4000/4000 [=====] - 63s 16ms/step - loss: 0.0053 - val\_loss: 0.0048  
Epoch 40/200  
4000/4000 [=====] - 61s 15ms/step - loss: 0.0052 - val\_loss: 0.0047  
Epoch 41/200  
4000/4000 [=====] - 62s 15ms/step - loss: 0.0051 - val\_loss: 0.0047  
Epoch 42/200  
4000/4000 [=====] - 61s 15ms/step - loss: 0.0051 - val\_loss: 0.0046  
Epoch 43/200  
4000/4000 [=====] - 62s 15ms/step - loss: 0.0050 - val\_loss: 0.0046  
Epoch 44/200  
4000/4000 [=====] - 61s 15ms/step - loss: 0.0049 - val\_loss: 0.0045  
Epoch 45/200  
4000/4000 [=====] - 62s 15ms/step - loss: 0.0049 - val\_loss: 0.0044  
Epoch 46/200  
4000/4000 [=====] - 62s 15ms/step - loss: 0.0048 - val\_loss: 0.0044  
Epoch 47/200  
4000/4000 [=====] - 61s 15ms/step - loss: 0.0048 - val\_loss: 0.0043  
Epoch 48/200  
4000/4000 [=====] - 61s 15ms/step - loss: 0.0047 - val\_loss: 0.0042  
Epoch 49/200  
4000/4000 [=====] - 62s 15ms/step - loss: 0.0046 - val\_loss: 0.0042

Epoch 50/200  
4000/4000 [=====] - 62s 15ms/step - loss: 0.0046 - val\_loss: 0.0042  
Epoch 51/200  
4000/4000 [=====] - 62s 15ms/step - loss: 0.0046 - val\_loss: 0.0041  
Epoch 52/200  
4000/4000 [=====] - 61s 15ms/step - loss: 0.0045 - val\_loss: 0.0041  
Epoch 53/200  
4000/4000 [=====] - 62s 15ms/step - loss: 0.0045 - val\_loss: 0.0040  
Epoch 54/200  
4000/4000 [=====] - 63s 16ms/step - loss: 0.0044 - val\_loss: 0.0040  
Epoch 55/200  
4000/4000 [=====] - 63s 16ms/step - loss: 0.0044 - val\_loss: 0.0039  
Epoch 56/200  
4000/4000 [=====] - 63s 16ms/step - loss: 0.0043 - val\_loss: 0.0040  
Epoch 57/200  
4000/4000 [=====] - 63s 16ms/step - loss: 0.0043 - val\_loss: 0.0038  
Epoch 58/200  
4000/4000 [=====] - 63s 16ms/step - loss: 0.0042 - val\_loss: 0.0038  
Epoch 59/200  
4000/4000 [=====] - 63s 16ms/step - loss: 0.0042 - val\_loss: 0.0038  
Epoch 60/200  
4000/4000 [=====] - 62s 16ms/step - loss: 0.0042 - val\_loss: 0.0037  
Epoch 61/200  
4000/4000 [=====] - 63s 16ms/step - loss: 0.0042 - val\_loss: 0.0037  
Epoch 62/200  
4000/4000 [=====] - 63s 16ms/step - loss: 0.0041 - val\_loss: 0.0037  
Epoch 63/200  
4000/4000 [=====] - 62s 15ms/step - loss: 0.0041 - val\_loss: 0.0037  
Epoch 64/200  
4000/4000 [=====] - 61s 15ms/step - loss: 0.0041 - val\_loss: 0.0036  
Epoch 65/200  
4000/4000 [=====] - 61s 15ms/step - loss: 0.0040 - val\_loss: 0.0035  
Epoch 66/200  
4000/4000 [=====] - 62s 15ms/step - loss: 0.0040 - val\_loss: 0.0035  
Epoch 67/200  
4000/4000 [=====] - 63s 16ms/step - loss: 0.0040 - val\_loss: 0.0036  
Epoch 68/200  
4000/4000 [=====] - 62s 15ms/step - loss: 0.0040 - val\_loss: 0.0034  
Epoch 69/200  
4000/4000 [=====] - 62s 16ms/step - loss: 0.0040 - val\_loss: 0.0036  
Epoch 70/200  
4000/4000 [=====] - 62s 16ms/step - loss: 0.0039 - val\_loss: 0.0034  
Epoch 71/200  
4000/4000 [=====] - 62s 15ms/step - loss: 0.0039 - val\_loss: 0.0034

Epoch 72/200  
4000/4000 [=====] - 62s 15ms/step - loss: 0.0039 - val\_loss: 0.0034  
Epoch 73/200  
4000/4000 [=====] - 62s 15ms/step - loss: 0.0038 - val\_loss: 0.0034  
Epoch 74/200  
4000/4000 [=====] - 62s 16ms/step - loss: 0.0039 - val\_loss: 0.0034  
Epoch 75/200  
4000/4000 [=====] - 61s 15ms/step - loss: 0.0038 - val\_loss: 0.0033  
Epoch 76/200  
4000/4000 [=====] - 61s 15ms/step - loss: 0.0038 - val\_loss: 0.0033  
Epoch 77/200  
4000/4000 [=====] - 61s 15ms/step - loss: 0.0038 - val\_loss: 0.0033  
Epoch 78/200  
4000/4000 [=====] - 62s 15ms/step - loss: 0.0038 - val\_loss: 0.0032  
Epoch 79/200  
4000/4000 [=====] - 62s 15ms/step - loss: 0.0037 - val\_loss: 0.0033  
Epoch 80/200  
4000/4000 [=====] - 61s 15ms/step - loss: 0.0037 - val\_loss: 0.0032  
Epoch 81/200  
4000/4000 [=====] - 61s 15ms/step - loss: 0.0037 - val\_loss: 0.0032  
Epoch 82/200  
4000/4000 [=====] - 62s 16ms/step - loss: 0.0037 - val\_loss: 0.0032  
Epoch 83/200  
4000/4000 [=====] - 62s 16ms/step - loss: 0.0037 - val\_loss: 0.0031  
Epoch 84/200  
4000/4000 [=====] - 62s 15ms/step - loss: 0.0036 - val\_loss: 0.0032  
Epoch 85/200  
4000/4000 [=====] - 63s 16ms/step - loss: 0.0036 - val\_loss: 0.0031  
Epoch 86/200  
4000/4000 [=====] - 62s 15ms/step - loss: 0.0036 - val\_loss: 0.0032  
Epoch 87/200  
4000/4000 [=====] - 62s 15ms/step - loss: 0.0036 - val\_loss: 0.0031  
Epoch 88/200  
4000/4000 [=====] - 62s 16ms/step - loss: 0.0036 - val\_loss: 0.0031  
Epoch 89/200  
4000/4000 [=====] - 61s 15ms/step - loss: 0.0036 - val\_loss: 0.0031  
Epoch 90/200  
4000/4000 [=====] - 62s 15ms/step - loss: 0.0035 - val\_loss: 0.0032  
Epoch 91/200  
4000/4000 [=====] - 62s 15ms/step - loss: 0.0035 - val\_loss: 0.0031  
Epoch 00091: early stopping

4000/4000 [=====] - 62s 15ms/step - loss: 0.0035 - val\_loss: 0.0031  
Epoch 00091: early stopping




**9. Presenta una tabla que compare el modelo seleccionado contra algunos de los otros modelos que hayas evaluado (los más prometedores). Dependiendo del problema que hayas decidido atacar, la comparación puede incluir tiempo de entrenamiento, pérdida final, error de clasificación.**

	Pérdida de validación (Epoch 10)	Diferencia entre pérdida de validación y de entrenamiento (sobre-ajuste)	Tiempo de entrenamiento (Epoch 5)	Número de parámetros entrenables
Original (Filtros 8 y tamaño de batch 512)	0.0201	0.0070	76s	138,601
Aumentando número de filtros	0.0172	0.0070	716s	1,698,057
Aumentando tamaño de batch (batch size = 32)	0.0087	0.0013	555s	1,698,057
Aumentando tamaño de batch (batch size = 64)	0.0072	0.0006	546s	1,698,057
Aumentando tamaño de batch (batch size =128)	0.0066	0.0006	546s	1,698,057
Aumentando tamaño de batch (batch size =256)	0.0064	0.0006	554s	1,698,057
Quitando capa densa	0.0054	0.0008	1937s	426,447
Quitando capa densa y capas duplicadas de convoluciones (aún hay 6 capas convolucionales)	0.0048	0.0007	1170s	186,441
Simplificar modelo pasado. Borrar los os capas convolucionales. Total 4 capas convolucionales.	0.0067	0.0002	271s	38,793

Simplificar modelo pasado reduciendo número de filtros (16,32,16)	0.0094	0.0013	397s	10,185
Simplificar modelo pasado reduciendo número de filtros aún más (8,16,8) [MODELO SELECCIONADO]	0.0121	0.0000	72s	2,793
Modelo pasado con 91 épocas	0.0031	0.0004	62s	2,793
Simplificar modelo pasado reduciendo número de capas convolucionales a 2.	0.0163	0.0087	69s	921
Simplificar modelo pasado reduciendo número de capas convolucionales a 1.	0.0593	0.0109	31s	96



**10. Presenta ejemplos de la reconstrucción o predicción de imágenes para los casos de denoising o segmentación, o la matriz de confusión para el caso de extracción de representaciones.**





## 11. Variación de ruido.

El original usa el 20% de ruido. A continuación, se muestran las imágenes originales con mayor porcentaje de ruido y su respectiva reconstrucción.



Salt & Pepper	MSE
20%	0.0031
30%	0.0040
40%	0.0052
50%	0.0073
60%	0.0107
80%	0.0275
90%	0.0447
100%	0.0594

Tabla mostrando el error cuadrático medio con mayores porcentajes de ruido